



LIVEQUEUE

We hold your spot. You own your time.

LiveQueue

A Real-Time Digital Queue Management Platform

Join from a QR code, wait anywhere, get called the moment it is your turn — with server-authoritative dispatch across many counters.

In production • Android v1.0.6

Presenter: [Your name] • Supervisor: [Name] • [Department, University] • [Date]

Queues are still physical and uncertain

1 Visitors

- Stand in a crowd for an unknown time
- No live update, easy to miss a call
- Multi-step visits restart the wait

2 Counter staff

- Who is next? Who was here first?
- Several counters, several services
- Hand-offs between counters are informal

3 Administrators

- No visibility of load or idle counters
- No record of who served whom
- Hard to run several desks or teams

Multi-service visits (e.g. an enquiry, then a payment) make every one of these worse.

One live queue, shared by every device

Organizations & roles

Head, Manager, Admin
workspaces, Executives

Queues & counters

Many counters per queue ·
Open / Paused / Off

QR entry

One organization QR →
choose a queue

Ordered journeys

Several services in a chosen
order · referrals

Live updates

Socket.io: position, ETA,
called, done

Notifications

Android FCM · iPhone Web
Push · email

Governance & audit

Admin replacement · Head
succession · immutable
history

Floating Counter Console

Always-visible serving
controls (live)

Clients: Flutter Android app · React staff dashboard · iPhone/iPad Safari portal (PWA)

Four roles: governance and operations

GOVERNANCE

Organization Head

Whole organization · exactly one

Organization Manager

Org-wide reports & audit · no serving

OPERATIONS — isolated Admin workspaces

Admin A

Owens ≤ 1 live queue

Executive

Executive

Queue: Student Services

Admin B

Owens ≤ 1 live queue

Executive

Executive


Queue: Library Desk

✗ no cross-workspace access

Product rules (enforced in the database)


- One Head per organization
- One live queue per Admin
- A live queue always keeps its Admin
- Executives belong to one workspace
- Permissions derive from role only

System architecture



REST for truth, Socket.io for change

REST — every read and write




Socket.io — "something changed"




On reconnect: re-join rooms and refetch everything — a missed event is never assumed to be replayed. **Rooms:** organization (Head/Manager), workspace (Admin/Executives), token (one visitor), public queue (no personal data).

Token lifecycle: one server-side state machine



Serve next and Start are atomic compare-and-swap updates; the server rejects any transition the machine does not allow.

Ordered journeys and referrals



Ordered, 1-20 steps

Chosen by the visitor from the recommended order

Repeat rules

A service may repeat (limit 1-10), never twice in a row

Locked at creation

Any change is refused: 409 JOURNEY_LOCKED

Counter eligibility

A counter serves only the services it handles

Referral

Next step goes to a chosen open counter — no interruption

Fallback

Target closes → priority kept at another eligible counter

Many counters, one server-side dispatcher

Counter states

Open

ACTIVE · operator required

Paused

ON_BREAK · operator retained

Off

OFFLINE · unassigned

Assign operator to Off → Paused (never straight to Open). Off is refused while someone is called or in service there.

Serve next — decided on the server

1

Client sends only the queue id

No person, counter or staff id is accepted

2

Lock the caller's own counter

SELECT ... FOR UPDATE inside a transaction

3

Pick: referrals first, then earliest arrival

Whose current step this counter handles

4

Compare-and-swap WAITING → CALLED

A concurrent claim matches nothing and moves on


5

Commit, then broadcast

Every device converges on the same truth


Floating Counter Console

Compact



Real screenshots, fictional demo data.


Expanded




- Document Picture-in-Picture: a small always-on-top window (Chromium desktop browsers)
- Same APIs, same Socket.io connection — no second dispatch engine
- Shows token, service, state, waiting count, server-timed timer
- Actions held while reconnecting; access loss closes it
- Token number and service only — no personal data
- Fallback: in-page dock that says it cannot stay on top

Live in production • 9 Oct 2026


Visitor experience: QR → services → token



1 • Scan organization QR



2 • Choose services in order



3 • Live token, position, ETA

Android (Flutter)

Same flow in the app; push via Firebase Cloud Messaging.

iPhone / iPad


Safari portal; notifications need iOS 16.4+, Home Screen install and permission.

Physical iPhone push: pending

Physical Android QA: pending

Screens: fictional demo data.

Notification architecture



Governance and succession



Admin replacement

Queue + every Executive move atomically; counters, tokens and history untouched



Live queue keeps its Admin

Database trigger; refused with a guided replacement flow



Head succession

Password + emailed code → single-use 72-hour link → explicit acceptance



Tenure history

Who led the organization, when, and why — visible by role



Transactional audit

Change and audit row commit together, with actor snapshots



Immutable history

Triggers block edits and deletes; deleting an org leaves a minimal receipt

Security: implemented vs known gaps


Implemented

- bcrypt passwords · 15-min JWT · rotated, hashed refresh tokens
- User + organization reloaded from the DB on every request
- Tenant & workspace scoping on every query (404, never leak)
- Backend authorization on every mutation, whatever the UI shows
- Instant revocation: tokens + live sockets dropped
- Single-use, hashed, expiring invitation and handover links
- Transactional, immutable audit with actor snapshots
- Rate limits per real client (CF-Connecting-IP; spoof-resistant)

Known gaps

- Rate limits are in-memory, per process
- People behind one NAT (campus Wi-Fi, carrier) share a limit
- No external penetration test or WAF
- No centralized security monitoring or alerting


Database design (simplified)



Enforced in PostgreSQL


- One Head per org
- One live queue per Admin
- Live queue needs an Admin
- One open succession
- Unique token number per queue
- History append-only (triggers)

Deployment and release pipeline



Single instance by design: Socket.io has no cross-instance adapter yet.

Release process (used for v1.0.6)



APK verification on every release: exactly one signer • certificate pinned to 70:08:D3:....0A:93 • package com.livequeue.mobile_app • version 1.0.6 (build 7) • not debuggable • production API compiled in. Signing keys live only in GitHub secrets — never in the repository.

Testing and current production status

1,234

backend tests

640

dashboard tests

479

Flutter tests

All passing on production master (9 Oct 2026). Typecheck, lint and production builds clean.

Verified automatically

- Backend live: health, Socket.io handshake, Web Push config
- Both governance migrations applied; new routes respond
- Dashboard + portal render (headless Chromium checks)
- Android v1.0.6 published, Latest; APK independently verified
- Console + performance fixes live; per-client rate limits verified

Pending / not yet done

- Real iPhone Web Push delivery — pending physical iPhone test
- Physical Android device QA — pending
- Formal load testing — not done
- Console always-on-top on a real desktop OS — not verified headlessly
- Post-deploy DB counts recorded manually by the owner (read-only SQL)

Capacity, cost and scaling

Answering “How big can it run for free?” — with every number labelled

What is actually free today?


Service	Plan	Key free limits (PROVIDER LIMIT)
Render (API + Socket.io)	Free web · Hobby	0.1 CPU · 512 MB · 1 instance · sleeps after 15 min (≈1 min wake) · 5 GB outbound/month
Neon (PostgreSQL)	Free	100 CU-h/month (≈400 h at 0.25 CU) · 1 GB storage · 5 GB egress · sleeps after 5 min
Cloudflare Pages	Free	500 builds/month · no published static bandwidth cap
Firebase Cloud Messaging	No-cost	600k messages/min per project
Web Push (VAPID)	No fee	iPhone/iPad Safari Home Screen PWA via Apple's push service; iOS/iPadOS 16.4+ and user permission
Resend	Free	100 emails/day · 3,000/month
GitHub	Free (public repo)	Standard Actions runners free for public repos

Domain (tdastudbook.au): separate registration/renewal cost — not verified. Quotas checked on: 2026-10-08 · production state: 2026-10-09

Registered ≠ daily active ≠ waiting now ≠ open sockets ≠ requests/second. Infrastructure is sized by concurrency and event rate, not by accounts.

Capacity: one queue

Socket bytes to one staff tab per serve cycle
(MEASURED, local)



ETA recompute per event (local): 16 ms · 103 ms · 454 ms

1 Serve next, 200 waiting: ~400 → 4 dashboard requests · 0 visitor re-reads
(MEASURED, local)

	People waiting in one queue
Application limit	None (HARD LIMIT: 32-bit token number)
Comfortable	~300 – 500 on free CPU (ESTIMATE)
Caution	> 500, or many open staff tabs (ESTIMATE)
Before the fixes	≤ ~50 (ESTIMATE)
Load-tested	NOT LOAD-TEST VALIDATED

Counters per queue: no limit; comfortable 1–10, caution 10–25 (ESTIMATE).

Capacity at every level

Level	Application rule (HARD LIMIT)	Comfortable (ESTIMATE)	Upgrade trigger
One counter	One operator · one person at a time	~12 visits/hour (5-min service)	—
One queue	No waiting limit	~300–500 waiting · 1–10 counters	> 500 waiting · > 25 counters
One organization	1 live queue per Admin; no Admin limit	1–2 queues · ~200 visits/day	≥ 5 busy queues
Whole platform	No organization limit	≈ 25,000 visits/month · ~5 small organizations	Neon free quotas exhausted

Assumptions: 4 queue events per visit, ~10 waiting, 2 staff tabs per queue, 8 h/day, 22 days/month. NOT LOAD-TEST VALIDATED.

Current bottlenecks, in order



Fixed and live (9 Oct 2026)

Dashboard refetch storm · iPhone portal re-reads · shared rate-limit bucket



Neon free quotas: 5 GB transfer + 100 CU-h/month

First likely limit; the reminder job keeps the DB awake while the backend is



Render Free: 0.1 CPU, ~1-minute cold start

Per-event ETA work grows with queue length



Render outbound: 5 GB/month

Socket updates grow with waiting × open staff tabs



Resend: 100 emails/day

Only if queues verify every visitor by email



Not limiting now: Cloudflare Pages, FCM, Web Push, GitHub Actions

Free or no published cap at this scale

ESTIMATE from measured payloads + published quotas. Quotas checked on: 2026-10-08 · production state: 2026-10-09


When do we need to pay?

When this happens	Service	Upgrade to	Why	Approx. cost
Refetch storm / shared limits	Code	Done — live 9 Oct 2026	Batched refetch · per-client limits	\$0
Neon egress or hours run out	Neon	Launch (usage-based)	500 GB egress, no hour cap	≈ \$11–20/mo*
Cold starts / CPU busy	Render	Starter 0.5 CPU	Always on, 5× CPU	\$7/mo†
Outbound > 5 GB	Render workspace	Payment method or Pro	Avoid suspension; 25 GB	\$0.15/GB or \$25/mo
Sustained load	Render	Standard 1 CPU · 2 GB	More compute, 1 instance	\$25/mo†
> 100 emails/day	Resend	Pro	50k/month, no daily cap	\$20/mo
Need > 1 instance	Render + Redis	Replicas + Socket.io adapter	Events reach every instance	Size after load test


* $0.25 \text{ CU} \times 400\text{--}730 \text{ h} \times \$0.106/\text{CU-h}$. † Render pricing page is dynamic; figures corroborated by third-party listings — confirm in dashboard. Quotas checked on: 2026-10-08 · production state: 2026-10-09

From one backend to distributed realtime

TODAY




FUTURE (not implemented)



Adding instances alone breaks realtime: events stay on the instance that emitted them.

Cost & capacity roadmap



ESTIMATES, not load-test validated. Fixed vs usage-based: Render plans fixed; Neon Launch usage-based. Domain cost separate. Quotas checked on: 2026-10-08 · production state: 2026-10-09

Known limitations and future development

Limitations today

- Single backend instance; no Socket.io adapter
- Free-tier cold start (~1 minute after 15 idle minutes)
- Per-person socket fan-out (grows with waiting × staff tabs)
- In-memory rate limits; a shared NAT shares one limit
- No formal load test; no central metrics/alerts
- Physical iPhone push and Android QA pending

Future (not implemented)

- Queue-level realtime updates instead of one per person
- Redis + Socket.io adapter + horizontal replicas
- Observability: metrics, central logs, alerts
- Analytics and predictive wait times
- Kiosk mode and institutional integrations
- Optional desktop (Tauri/Electron) counter companion

Live demo plan (≈ 5 minutes)

- 1 Log in as Head: dashboard, queue, counters, services
- 2 Visitor scans the organization QR on a phone
- 3 Build a 2-step journey; get a token
- 4 Executive presses Serve next — phone updates live
- 5 Start → Complete step → person returns for step 2
- 6 Referral to another counter (if time allows)
- 7 Open the Floating Counter Console (local build)
- 8 Show Associates, audit log and service history

Fallback: wake Render 5 minutes early · keep a local build ready · use the screenshots in this deck if a provider is down.

Conclusion

Real problem	Physical, uncertain queues → live, multi-counter, multi-service
Engineering	Server-authoritative dispatch · transactions · DB-enforced rules
Realtime	Socket.io across Android, iPhone and the dashboard
Trust	Instant revocation · immutable audit · governance & succession
Honest scale	Free today; first limit is egress (fixable); clear upgrade path

Questions?